



# Mastering PostgreSQL in Kubernetes with CloudNativePG

Course introduction

**Neel Patel** - Principal Software Engineer at EDB

**Danish Khan** - Senior SDE at EDB

March 5, 2025 - PGConf India, 2025

# About Danish

Senior SDE, EDB

- Certified Kubernetes application Developer
- Certified Kubernetes Admin
- DoK Ambassador
- Vmware certified VCP-DCP 2014
- Virtualization Enthusiast
- Open Source CloudNativePG contributer
- Major Contributor in patented unified test framework  
<https://patents.justia.com/patent/11366747>
- With EDB since ~5years



# About Neel

Principal Software Engineer at EDB

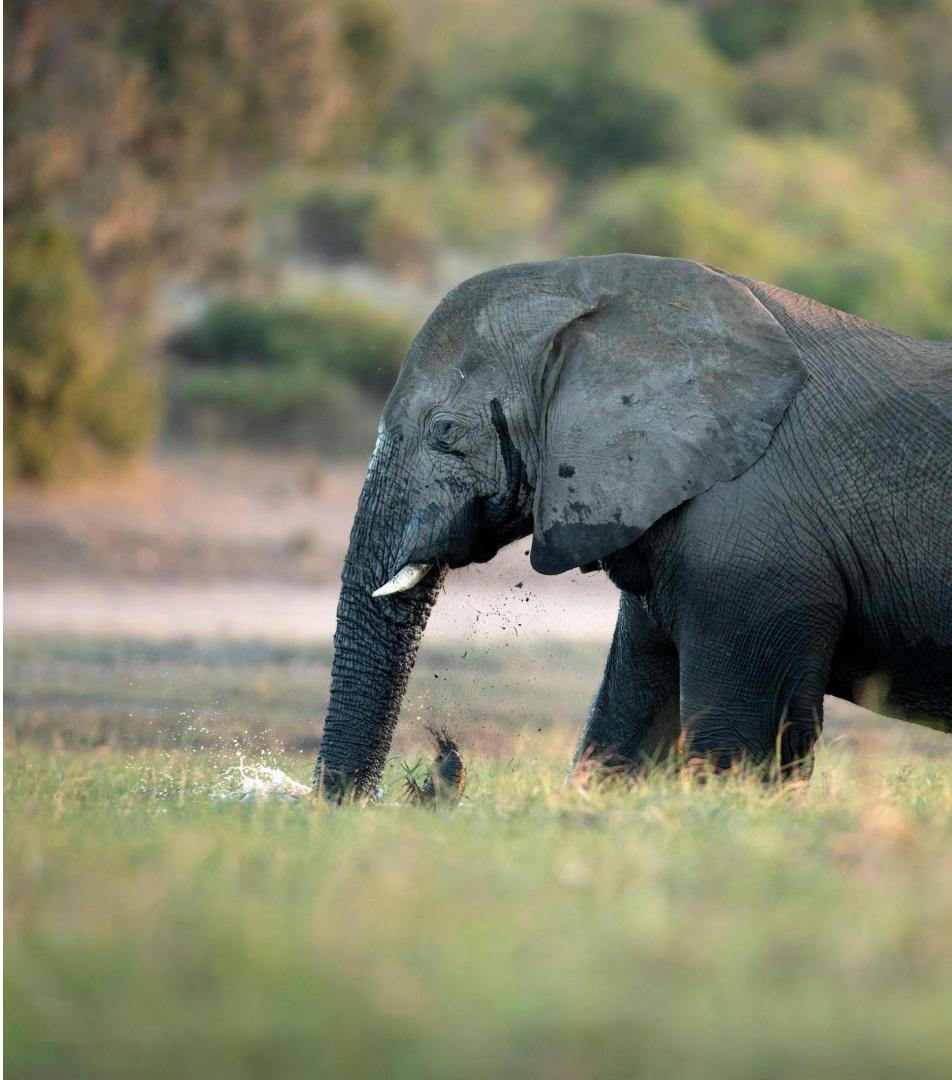
- With EDB since ~10 years
- Using PostgreSQL since ~ 10 years
- Open source contributor
- Kubernetes and GoLang Developer



# Agenda

- Overview of the open-source stack: PostgreSQL, Kubernetes, CloudNativePG  
Hands-on: Setting up the local environment and presenting the lab exercise (working in 2-3 person groups is encouraged). PostgreSQL 17 will be recommended for usage.
- Hands-on: Bootstrapping a new PostgreSQL cluster
- Hands-on: Configuring PostgreSQL
- Hands-on: Scaling up/down
- Hands-on: Connecting from a sample application
- Hands-on: simulating an automated failover
- Hands-on: Setting up continuous backup on object stores
- Hands-on: `cnpg` plugin to observe the status and on Demand Backup
- Hands-on: Full and Point-In-Time Recovery
- Metrics and Logging
- Hibernation & Rehydration
- Q&A





## Part 1

# PostgreSQL

# What is PostgreSQL?

The world's most advanced database. Also known as **Postgres**. URL: postgresql.org

- **100% Open Source**
  - Widely used, extremely robust, and feature-rich.
- **Extensible and Customizable**
  - Support for custom data types, functions, and procedural languages.
- **Advanced features**
  - Includes replication, partitioning, full-text search, and JSON support.
- **ACID-Compliant**
  - Ensures Atomicity, Consistency, Isolation, and Durability for transactions.
  - Trusted choice for mission-critical applications
- **Strong Community Support**
  - Backed by a large, active global community.





## Part 2

## Kubernetes

# What is Kubernetes?

Greek for helmsman. Also known as **k8s**. URL: kubernetes.io

- Orchestration system for containers
- Automates deployment, administration and scaling of:
  - **infrastructure**
  - **cloud native applications** - also known as **workloads**
- Open Source (Apache Licence 2.0) since 2014
  - v1.0 released in 2015
- Owned by the Cloud Native Computing Foundation (CNCF, [cncf.io](https://cncf.io))
- Written in Go



CLOUD NATIVE  
COMPUTING FOUNDATION



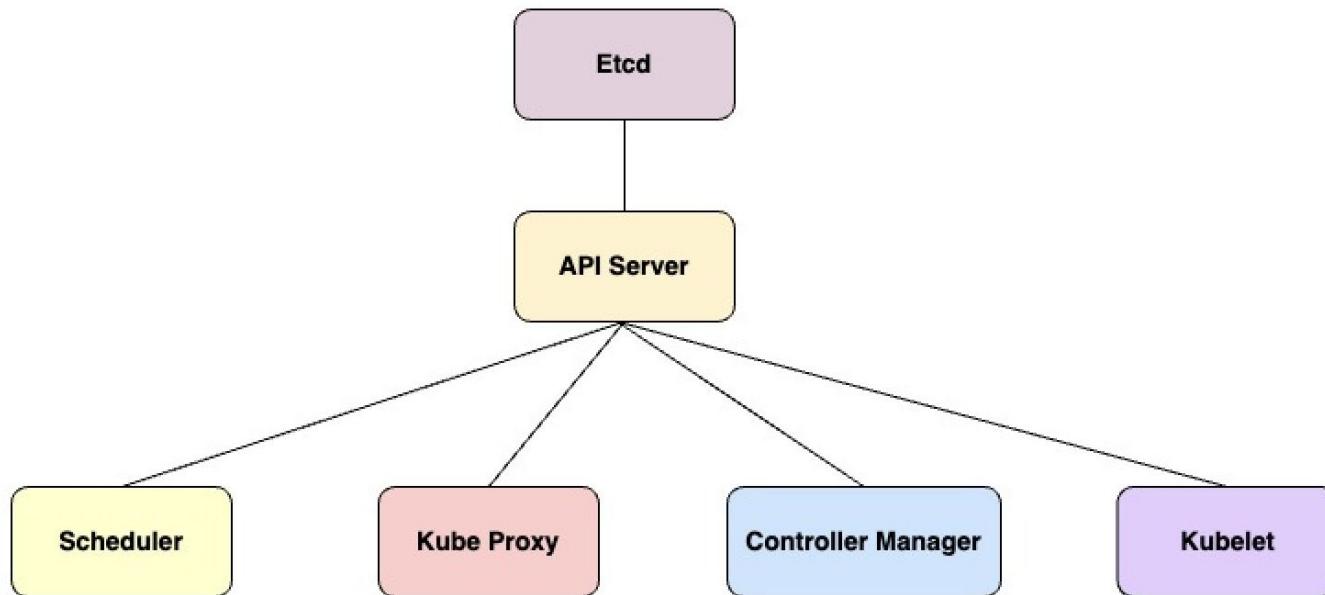


# Why Kubernetes?

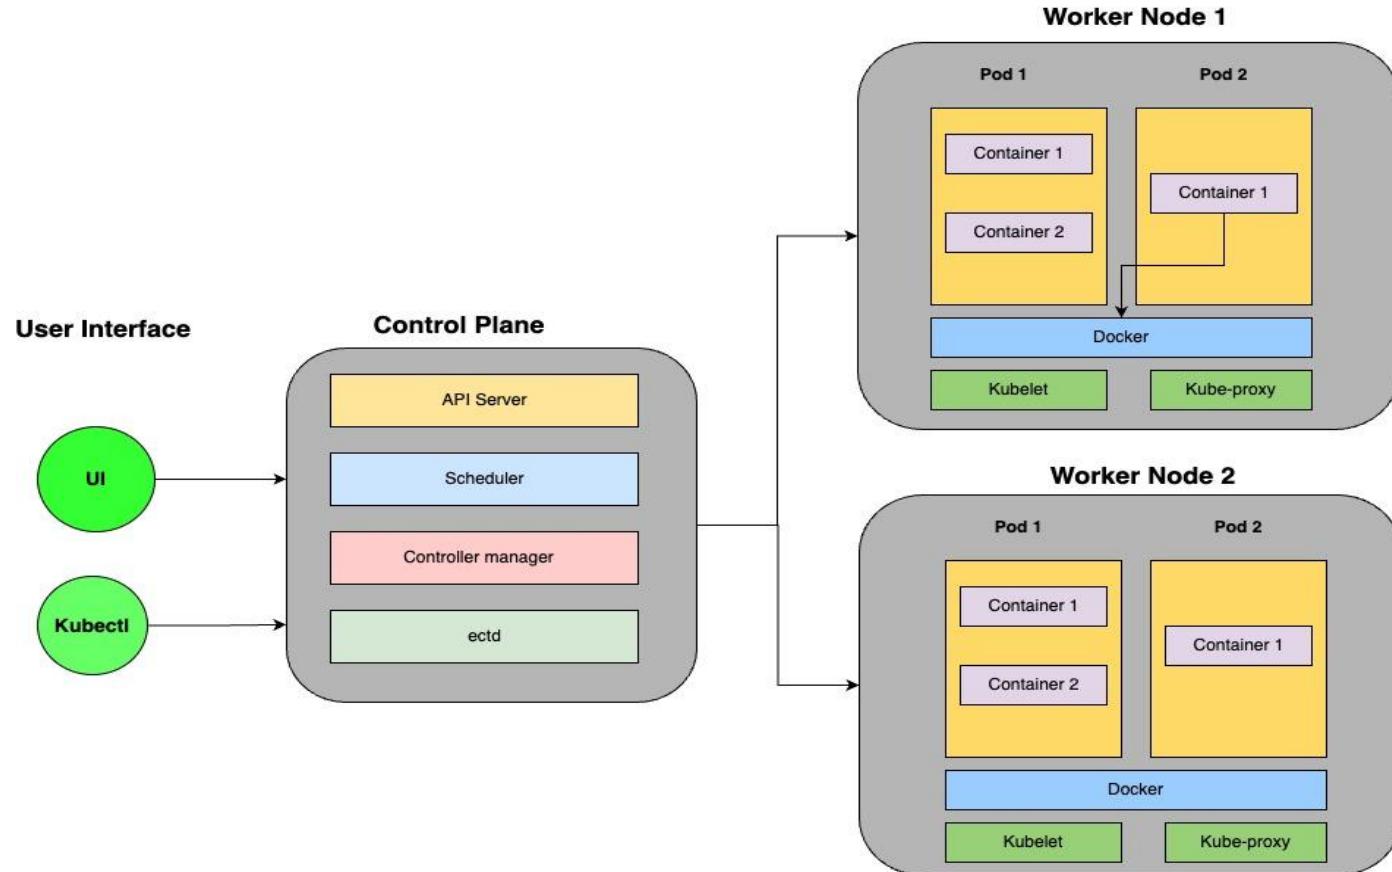
- **Flexibility:** Kubernetes works with many different infrastructures, including on-premise services, private and public clouds.
- **Efficiency:** Kubernetes optimizes the use of hardware resources, saving costs.
- **Self-healing mechanisms:** Kubernetes constantly checks the health of nodes and containers and restart/replace the failed node/pods if it is not responding to user-defined health checks.
- **Security and resource management:** Provides security features and efficient management of resources, ensuring that the infrastructure is secure and the resources are used optimally.
- Kubernetes helps developers reduce time to market by automating many of the manual tasks associated with application development and deployment.



# Main Components of Kubernetes



# Kubernetes Architecture



# Kubernetes vs Virtual Machines

## Kubernetes

Infrastructure and containerised apps

Containers share the same Kernel

Immutable Application Containers

Portable on any cloud, standard

## VMs

Operating systems

VMs have their own OS

Mutable (“dnf update”)

Indirectly portable



# From infrastructure to applications

Fundamental concepts for DBAs to better understand Kubernetes

- **Region (Kubernetes Cluster)**
- **Data Centre (Availability Zone)**
- **Nodes:**
  - Linux with K8s
  - VMs or Bare Metal
- **Network**
  - The Service resource
    - ClusterIP type
    - NodePort type
    - LoadBalancer type
- **Storage:**
  - Network vs Local
  - Local storage - Uses local disks attached to Kubernetes nodes.
  - Container Storage Interface (CSI)
  - The StorageClass Resource
- **Applications (workloads):**
  - The Deployment resource
  - The Statefulset resource
  - Databases are applications





## Part 3

# Kubernetes Operators

# What is Kubernetes operator ?

Extends kubernetes controller and it defines how complex applications work

- A Kubernetes operator automates actions of human being in a programmatic way
- A PostgreSQL database is a complex application:
  - Deployment and configuration
  - Failure detection and Failover
  - Updates and switchovers
  - Backup and Recovery
- Relies on Kubernetes' native components and its capabilities.
  - Self healing. High availability, scalability, resource control etc.
  - Declarative and fully managed



# Imperative vs Declarative with an example

## Imperative (VMs)

1. Create a PostgreSQL 17 instance
2. Configure for replication
3. Clone a second one
4. Set it as a replica
5. Clone a third one
6. Set it as a replica

## Declarative (Kubernetes)

There's a PostgreSQL 17 cluster with 3 instances (i.e. one primary and two replicas). **At any time.**



# Extending the Kubernetes API

- Kubernetes API is extensible (similarly to PostgreSQL)
  - You can define new data types with custom resources
  - You can define a custom controller for a custom resource (or more)
    - You declare the state of the custom resource
    - The controller ensures that the **current** state matches the **declared** one
      - Foundation of self-healing
- The operator pattern is a development pattern
  - Designed to manage complex applications
  - Custom resources and custom controllers
  - Ensures a declarative API



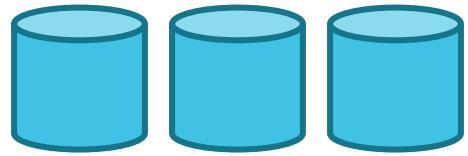
# The Simplified Kubernetes Resource Model

- Resources are identified by:
  - Namespace
  - Name
  - **Kind**/Type (e.g., Pod, Service)
- Resource Structure:
  - **spec** (Desired state)
  - **status** (Observed state)
- Key API Operations (Primitives):
  - **Get**: Retrieve resource information
  - **Put**: Modify a resource
  - **Watch**: Monitor resource changes over time



# A PostgreSQL 17 cluster resource with 3 instances

**spec (Desired State)**



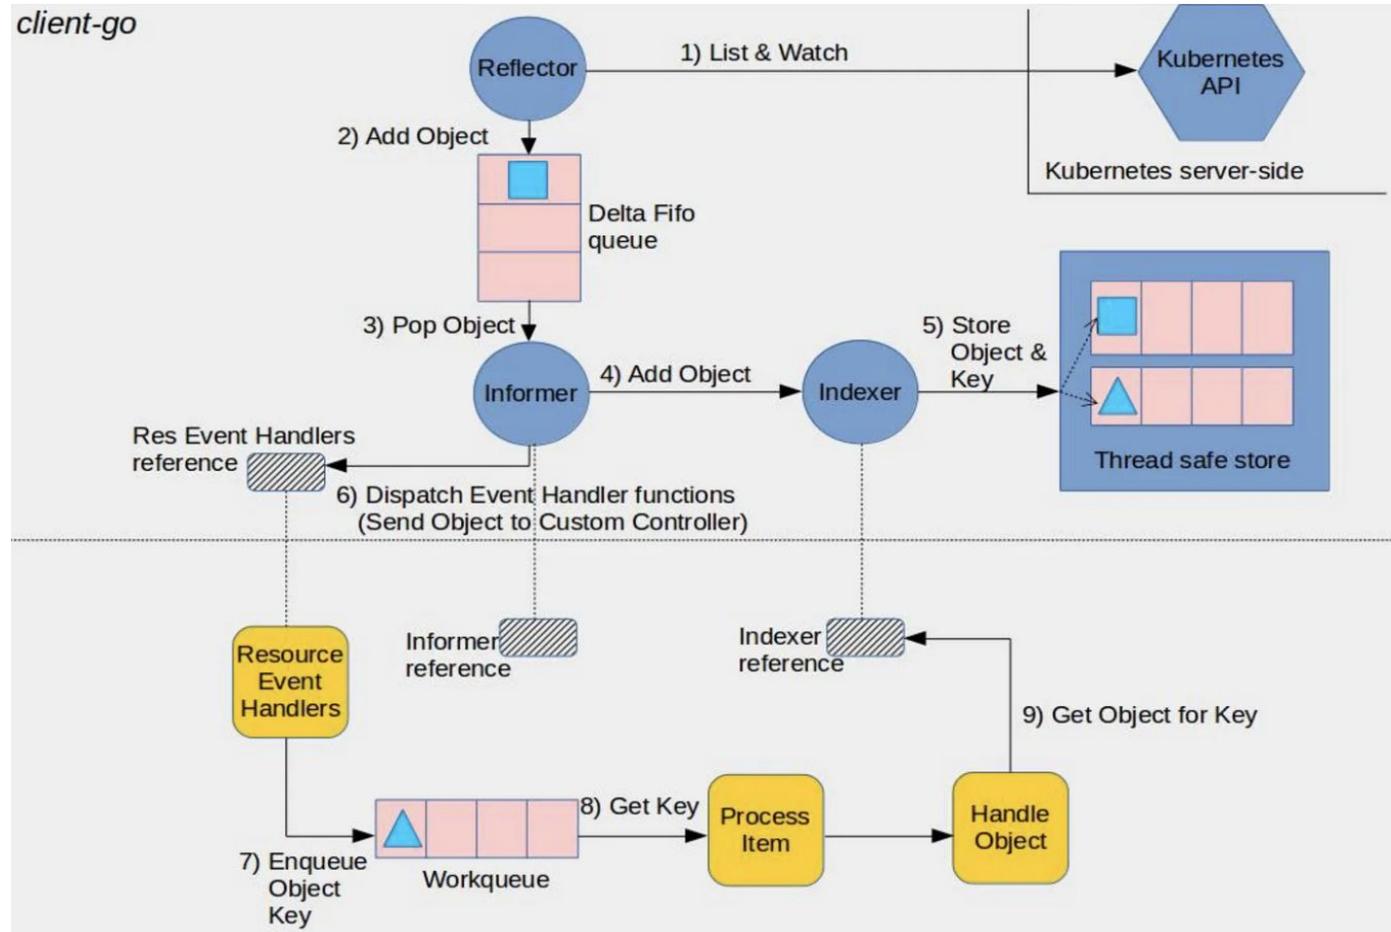
**status (Observed State)**

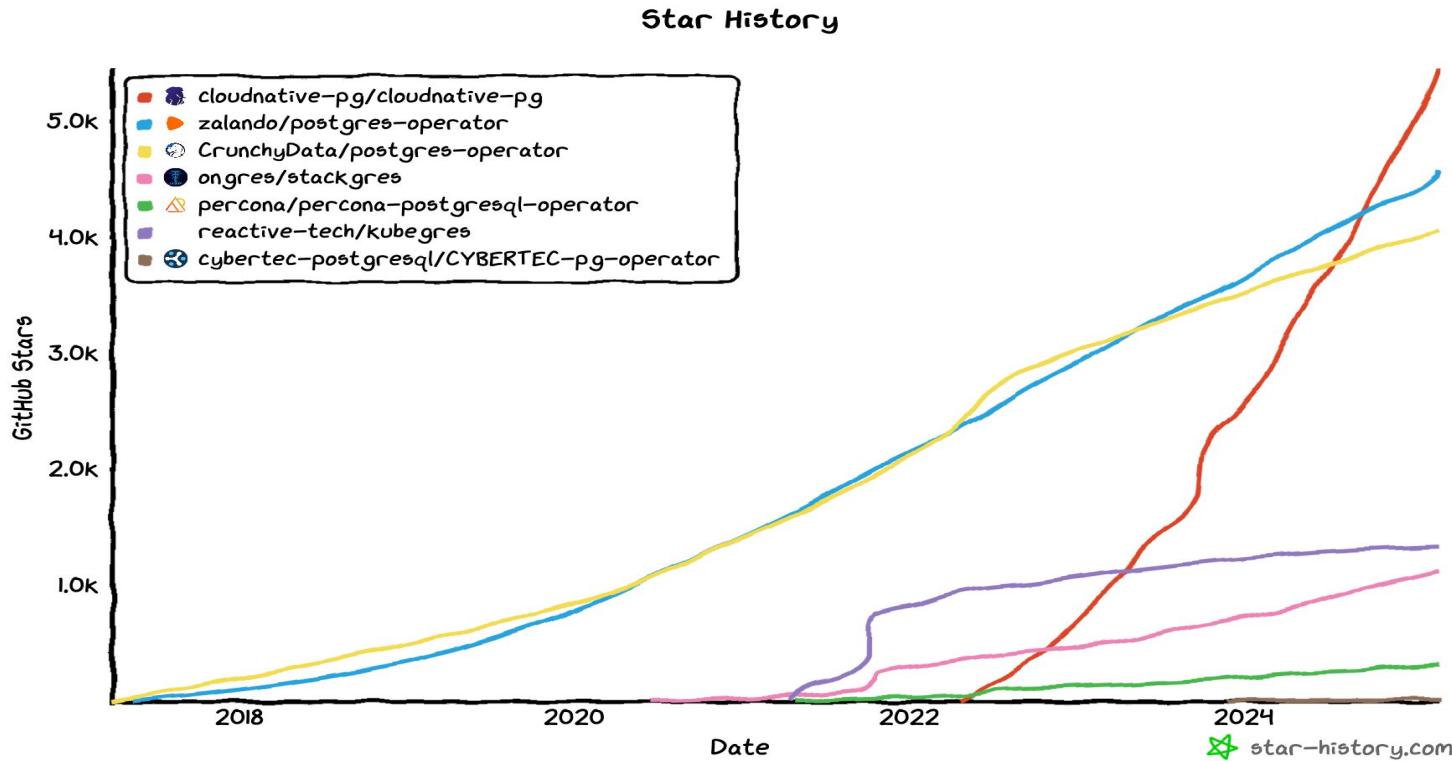
Nothing happens without a controller  
for the cluster resource

(simplified view)



# Reconciliation loop





CloudNativePG has reached the #1 spot for the first time, surpassing Zalando's PostgreSQL operator—a major milestone that highlights the growing impact and recognition of the project within the cloud-native ecosystem.





## Part 4

# CloudNativePG

# What is CloudNativePG?

“CloudNativePG is an open source operator designed to manage PostgreSQL workloads on any supported Kubernetes cluster running in private, public, hybrid, or multi-cloud environments. CloudNativePG adheres to DevOps principles and concepts such as declarative configuration and immutable infrastructure.”



# The history of CloudNativePG

**Bringing PostgreSQL to Kubernetes** (not Kubernetes to PostgreSQL)

## 2019-2020: The 2ndQuadrant era

- MVP (fail-fast)
- Cloud Native BDR (closed source)
- Cloud Native PostgreSQL (closed source)

## 2022-: CloudNativePG Revolution

- Launched in May 2022
- Consistent commitment from EDB
  - 10 releases in ~2 years
  - 3.4k commits
  - 120+ contributors, 300 forks
- Major features:
  - Declarative Fencing (2022)
  - Volume for WALs
  - Failover of physical Replication Slots
  - Backup from a Standby (2023)
  - Declarative Hibernation
  - Declarative Role Management
  - Volumes for Tablespaces
  - Volume Snapshot backup and recovery
  - Distributed topologies (2024)
  - Managed services

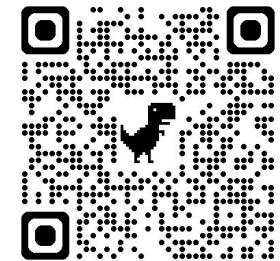
## 2021-2022: EDB Cloud Native PostgreSQL

- Closed source
- Production ready (IBM Cloud Pak)
- Several features gradually introduced
  - HA + DR
  - Observability
  - Security
  - Pooling



# About CloudNativePG

- “Level 5”, Production ready
  - EDB , IBM Cloud Pak, Google Cloud, Azure, Tembo, ...
- Open source (May 2022)
  - Apache License 2.0
  - Vendor neutral, openly governed, always free
  - Originally created by EDB (2019 by 2ndQuadrant)
- Multiple installation methods:
  - K8s manifests
  - Helm chart
  - OperatorHub.io (OLM)
- >59M downloads on Github (~2 years)
- Now CloudNativePG is Officially CNCF sandbox project



[github.com/cloudnative-pg](https://github.com/cloudnative-pg)



# The CloudNativePG Ecosystem

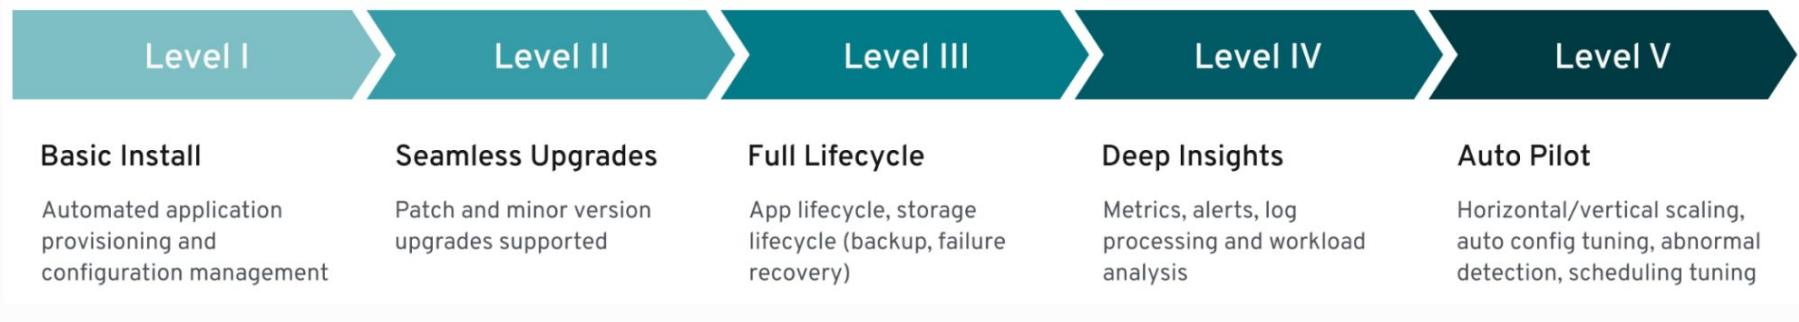
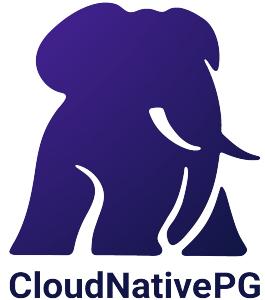
- The CloudNativePG operator is the main project (component)
  - 7 Maintainers (6 from EDB, 1 from Upbound)
  - Multiple component owners (contributors)
- Several components, with their respective owners:
  - Operand container images:
    - PostgreSQL with pgvector and pgaudit
    - PostgreSQL with PostGIS
    - PostgreSQL trunk (what will be in PostgreSQL 18 next year)
    - PgBouncer
  - Charts (operator and clusters)
  - Grafana Dashboard
  - Website
  - CNPG Playground
  - CNPG-I and official plugins (WIP)



[github.com/cloudnative-pg](https://github.com/cloudnative-pg)



# CloudNativePG capability levels



# The PostgreSQL `Cluster` resource

```
apiVersion: postgresql.cnpq.io/v1
kind: Cluster
metadata:
  name: clapton
spec:
  instances: 3
  imageName: ghcr.io/cloudnative-pg/postgresql:17.2
  affinity:
    nodeSelector:
      workload: postgres
  storage:
    size: 40Gi
  walStorage:
    size: 10Gi
```



# (Physical) Continuous backup

Guaranteed RPO from the first available backup to the latest archived WAL (not a snapshot)

## **WAL archive**

- Currently only via object store (Barman Cloud)
- WAL files are archived max every 5 minutes (RPO <= 5 min)

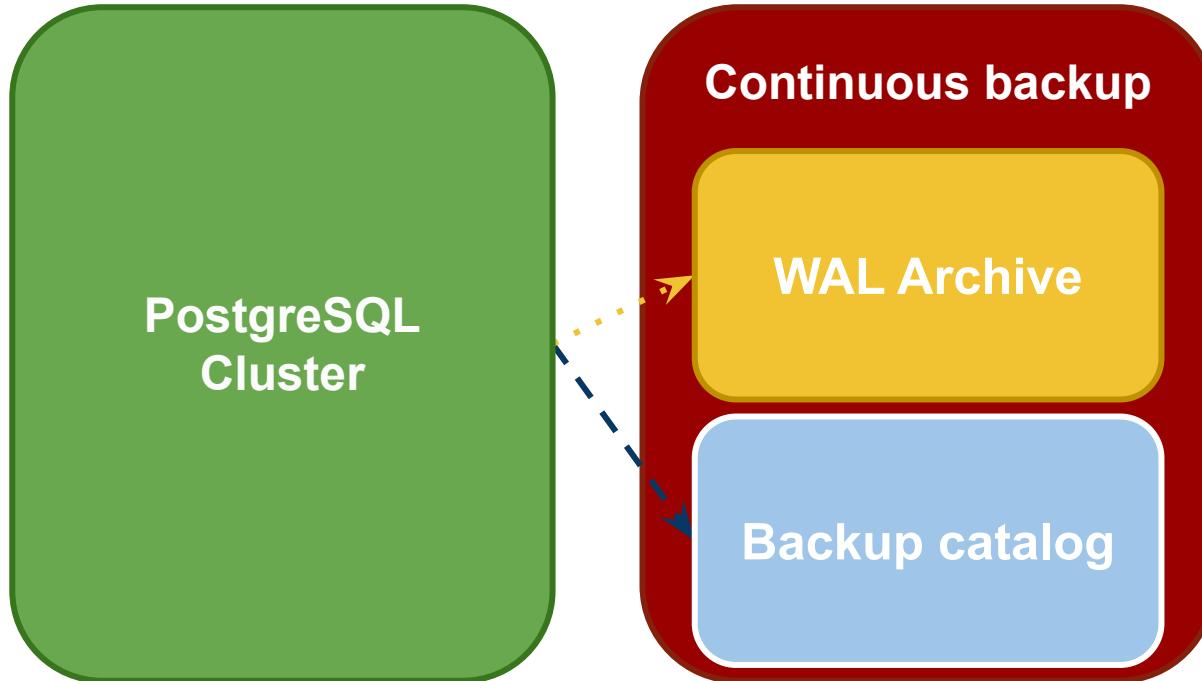
## **Physical base backups**

- From the primary or a standby
- Scheduled or on demand
- On object stores (Barman Cloud)
  - Hot (online)
- As Kubernetes Volume Snapshots
  - Hot (online) and cold (offline)
  - Transparent incremental/differential backup

## **Future: generic plug-in interface (CNPG-I)**



# Continuous backup



# Recovery

Full or to a specific point

## Bootstrap method

- Copy from an available physical base backup
- Apply WAL files (REDO logs) until you reach the target

## Full recovery

- Without a target (until the end of the WAL)
- Can be used to setup a physical replica
  - Replica cluster, even delayed
  - “Continuous recovery”

## Point-In-Time Recovery

- Up to a given time or transaction
- When reaching the target, the cluster promotes itself



# RTO vs RPO

- RPO is the Recovery Point Objective:

the maximum amount of data (measured in terms of time) that can be lost in case of a disaster.

In PostgreSQL clusters with transaction log archiving and point-in-time recovery, this is a function of how frequently WALs are shipped to the backup archive.

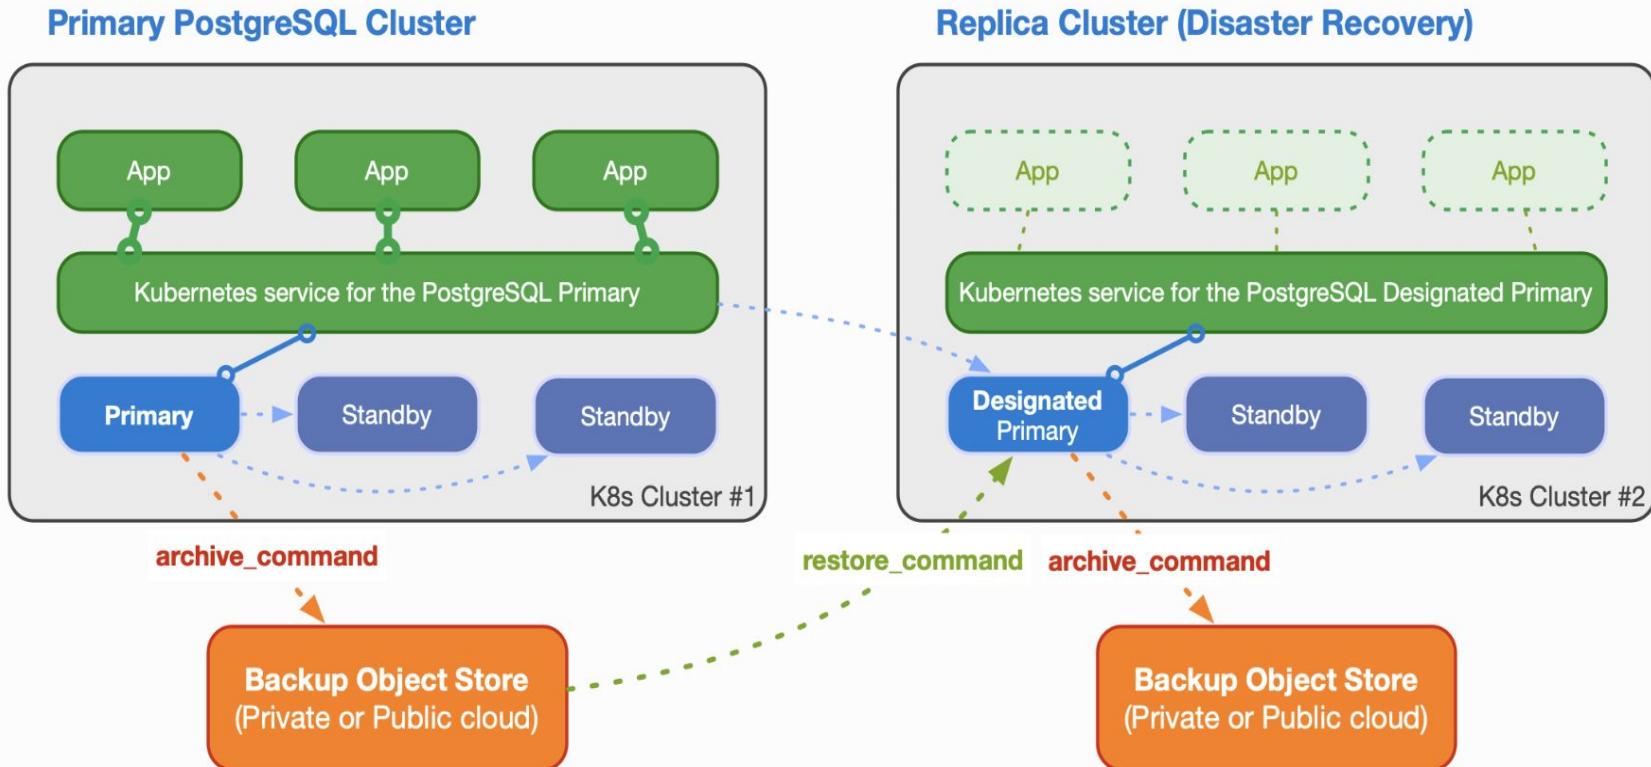
By default, Cloud Native PostgreSQL is configured to close a WAL file after a maximum of 5 minutes (“archive\_timeout” options).

- RTO is the Recovery Time Objective:

the maximum amount of time acceptable for a restoration to take. RTO can be extremely complicated, depending on the nature / scope of a disaster, hardware availability, snapshot frequency, etc.



# Replica cluster



# Storage

## For PostgreSQL data files

- Requires a volume with a file system on top of it, mounted inside a pod
- PVC with Filesystem volume mode
- PVC with ReadWriteOnce access mode

## For Backup data files

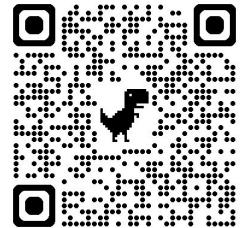
- WAL archive requires object stores
  - Future: PVCs (CNPG-I plugin)
- Backup catalog
  - Object stores
  - VolumeSnapshots
  - Hybrid (both object stores and VolumeSnapshots)
  - Future: PVCs (CNPG-I plugin)



# Direct management of Persistent Volume Claims (PVCs)

CloudNativePG is one of the few database operators that doesn't use Statefulsets

- Custom controller that directly manages PVCs instead of using Statefulsets
- Why?
  - `PGDATA` (database files) is the heart of a PostgreSQL database
  - Volumes for PGDATA, WAL files and tablespaces
    - They need to be coherent (PVC Group concept)
  - Online Resizing of volumes
  - Understanding of the primary/standby architecture in PostgreSQL
    - Rolling updates
    - Kubernetes upgrades
    - Management of standby sensitive parameters (e.g. max\_connections)
  - Advanced features made available by CSI for PVCs:
    - Online Resizing
    - PVC Cloning
    - Volume Snapshot backup and recovery
    - Volume Group Snapshot backup and recovery (alpha in K8s, with our active contribution)



# Monitoring in CloudNativePG

- Metrics
  - Native prometheus exporter for the operator
  - Native prometheus exported for each postgres instance
- Logs
  - Native support for the kubernetes events
  - Direct stdout logging of operator and PostgreSQL logs in JSON format



# Offline major version upgrades

- Offline major version upgrade
  - Leveraging logical backup methods like pg\_dump and pg\_restore

Limitations:

- Does not capture changes made to the origin database once the process begins
- Halt write operations on the origin's primary database until the migration is complete

Readiness:

- Verify that the import or upgrade processes function correctly on a new cluster, ensuring compatibility with applications on the updated Postgres version.

Tolerance:

- Accurately estimate the cutover time, primarily based on import duration



# Online major version upgrade using CloudNativePG

- Online major version upgrade
  - Leveraging PostgreSQL's native logical replication to minimize cutover downtime to nearly zero.

Limitations of Logical replication:

- Replication of Data Definition Language (DDL), such as *CREATE TABLE* and *ALTER TABLE*, does not occur.
- Sequences are not replicated.
- Large objects (LOBs) are excluded from replication.

With CloudnativePG:

- Managing database schema changes, as DDL is not replicated.
- Avoid schema modifications during upgrades or apply changes manually if needed.
- Updating sequences as part of the cutover process.



# How to achieve online upgrade using CloudNativePG

```
apiVersion: postgresql.cnpq.io/v1
kind: Cluster
metadata:
  name: pg13
spec:
  imageName: ghcr.io/cloudnative-pg/postgresql:13
  enableSuperuserAccess: false
  instances: 1
  storage:
    size: 1Gi
  bootstrap:
    initdb:
      dataChecksums: true
      walSegmentSize: 32
      postInitApplicationSQL:
        - CREATE TABLE songs (id SERIAL PRIMARY KEY, title TEXT)
        - ALTER TABLE songs OWNER TO app
        - INSERT INTO songs(title) VALUES ('Back in black')
        - CREATE PUBLICATION migrate FOR TABLE songs
  managed:
    roles:
      - name: user1
        ensure: present
        comment: User for logical replication connections
        login: true
        replication: true
        inRoles:
          - app
        passwordSecret:
          name: user1
  postgresql:
    pg_hba:
      - hostssl replication user1 10.0.0.0/8 md5
```

```
apiVersion: postgresql.cnpq.io/v1
kind: Cluster
metadata:
  name: pg16
spec:
  imageName: ghcr.io/cloudnative-pg/postgresql:16
  bootstrap:
    initdb:
      import:
        schemaOnly: true
        type: microservice
        databases:
          - app
        source:
          externalCluster: pg13
        postImportApplicationSQL:
          - CREATE SUBSCRIPTION migrate CONNECTION 'host=pg13-rw
user=user1 dbname=app sslmode=require password=user1pwd'
        PUBLICATION migrate
        externalClusters:
          - name: pg13
        connectionParameters:
          # Use the correct IP or host name for the source database
          host: pg13-rw
          user: user1
          dbname: postgres
          password:
            name: user1
            key: password
```



# Upgrade: Fully declarative way using CloudNativePG

Draft PR for community to validate:

<https://github.com/cloudnative-pg/cloudnative-pg/pull/6664>

What user have to do:

- Users can specify a new image with a higher major version

Process of upgrade:

- Hibernating the cluster which will delete cluster pod to ensure no write to data directory
- Create upgrade job on primary instance
- Job will use "*pg\_upgrade --link*" to efficiently migrate the data
- Once the migration is complete, all non-primary instances will be removed and new standbys are cloned as per the desired instance.



# Thank you

All the Maintainers, contributors, users and community members of CloudNativePG

**Want to contribute?**

<https://github.com/cloudnative-pg/cloudnative-pg/blob/main/CONTRIBUTING.md>

**EDB**

For generously donating the intellectual property for this project and for their decades of invaluable contributions to the broader PostgreSQL ecosystem.



# Time for Practice



# What are the requirements for the setup?

Laptop: Sufficient RAM, CPU, and disk to run Kubernetes in Docker (Kind) and any operating system that supports it.

Software: Students must come to the training session with the following software already installed, running the latest available version at the time of training:

- Kind (please follow instructions at [kind.sigs.k8s.io/docs/user/quick-start](https://kind.sigs.k8s.io/docs/user/quick-start))
- Kubectl
- The cnpg plugin for Kubectl (<https://cloudnative-pg.io/documentation/1.25/kubectl-plugin/>)

## Additional requirements

- curl
- jq
- stern
- kubectx
- Kubectl plugins:
  - view-secret
  - view-cert



# Exercise #1 CloudNativePG Installation

- Create kind cluster using ( kind should be installed in local machine )  
*"kind create cluster --name cnpg"*  
// get kind clusters  
**kind get clusters**
- Deploy the operator and ensure it is installed on the control plane node.  
**kubectl apply --server-side -f**

<https://raw.githubusercontent.com/cloudnative-pg/cloudnative-pg/release-1.25/releases/cnpg-1.2.5.0.yaml>

- Make sure CNP operator gets installed using  
**kubectl get deployment -n cnpg-system cnpg-controller-manager**
- Install CNPG kubectl plugin  
**curl -sSfL \**  
*<https://github.com/cloudnative-pg/cloudnative-pg/raw/main/hack/install-cnpg-plugin.sh> | \*  
**sudo sh -s -- -b /usr/local/bin**
- Validate version of CNPG kubectl plugin installed  
**kubectl cnpg version**



## Exercise #2 Bootstrap new PostgreSQL cluster

Open the “Quickstart” page from the documentation and download the `cluster-example.yaml` file from the website and save it locally. For example, you can use `curl`:

```
$ curl -LO  
https://raw.githubusercontent.com/cloudnative-pg/cloudnative-pg/refs/heads/main/docs/src/sam-  
ples/cluster-example.yaml
```

Open the `cluster-example.yaml` file and inspect it.

```
$ kubectl apply -f cluster-example.yaml
```

Get the YAML definition from K8s.

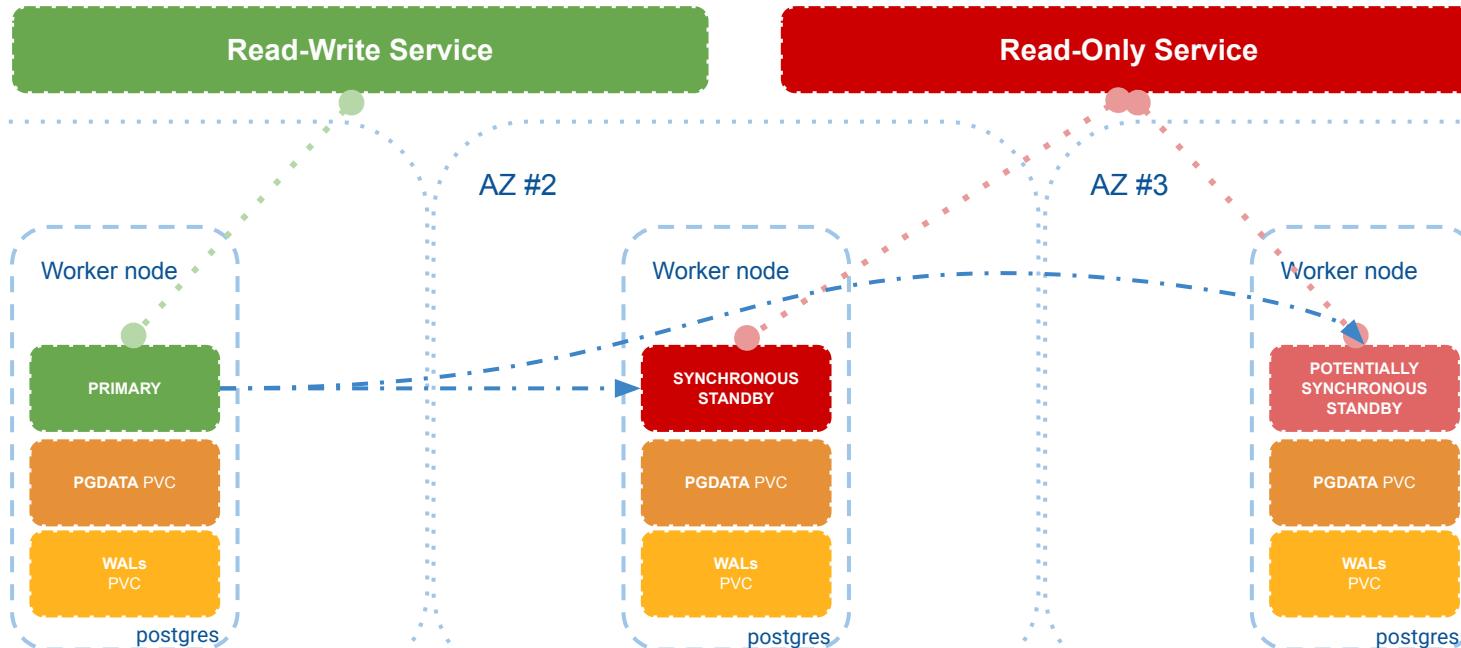
```
$ kubectl get cluster -o yaml
```

It will be different from the one we applied. Why?



# Highly Available PostgreSQL Cluster

K8s cluster



## Exercise #3 **Explore generated kubernetes resources**

```
$ kubectl get cluster
```

```
$ kubectl get pod
```

```
$ kubectl get service
```

```
$ kubectl get pvc
```

```
$ kubectl get pdb
```

```
# Optional
```

```
$ kubectl get role
```

```
$ kubectl get rolebinding
```

```
$ kubectl get sa
```



## Contd....

```
// inspect the users
```

```
kubectl exec -ti -c postgres cluster-example-1 -- psql -c '\du'
```

```
// inspect the databases
```

```
kubectl exec -ti -c postgres cluster-example-1 -- psql -c '\l'
```



# Contd....

```
$ kubectl get secrets -A
```

NAME	TYPE
cluster-example-app	kubernetes.io/basic-auth
cluster-example-ca	Opaque
cluster-example-replication	kubernetes.io/tls
cluster-example-server	kubernetes.io/tls

**cluster-example-app**: secret containing information about the app user in the PostgreSQL cluster, including credentials

**cluster-example-ca**: secret containing the Certification Authority (CA) used by the operator for emitting TLS certificates

**cluster-example-replication**: secret with the TLS client certificate for streaming replication in the HA cluster

**cluster-example-server**: secret with the TLS certificate of the PostgreSQL server



# Contd....

## Inspect CA certificate

```
$ kubectl get secret cluster-example-ca -o jsonpath=".data['ca\.crt']" | base64 -d | openssl x509 -text -noout
```

## Server certificate signed with above CA

```
$ kubectl get secret cluster-example-server -o jsonpath=".data['tls\.crt']" | base64 -d | openssl x509 -text -noout
```

## Certificate for the streaming\_replica user

```
$ kubectl get secret cluster-example-replication -o jsonpath=".data['tls\.crt']" | base64 -d | openssl x509 -text -noout
```

## Check, how standby is connecting to the primary cluster

```
$ kubectl exec -ti -c postgres cluster-example-2 -- psql -qAt -c 'SHOW primary_conninfo'
```

```
host=cluster-example-rw
user=streaming_replica
port=5432
sslkey=/controller/certificates/streaming_replica.key
sslcert=/controller/certificates/streaming_replica.crt
sslrootcert=/controller/certificates/server-ca.crt
application_name=cluster-example-2
sslmode=verify-ca
```



## Exercise #4 CNPG CRD

- Try changing the image from [ghcr.io/cloudnative-pg/postgresql:16.1](https://ghcr.io/cloudnative-pg/postgresql:16.1) to: ([ghcr.io/cloudnative-pg/postgresql:15](https://ghcr.io/cloudnative-pg/postgresql:15)).

Is that possible? Why?

- Try changing the ***shared\_buffers*** parameters to “thistest”.

Is that possible? Why?



# Exercise #5 Manual PostgreSQL administration activities

```
$ kubectl log cluster-example-1 | jq
```

```
$ kubectl cnpg psql cluster-example
```

```
psql (16.1 (Debian 16.1-1.pgdg110+1), server 16.2 (Debian 16.2-1.pgdg110+2))
```

```
Type "help" for help.
```

```
postgres=# \! hostname
```

```
cluster-example-1
```

```
$ kubectl cnpg psql --replica cluster-example
```



# Exercise #6 Testing the controller

Try deleting the -rw service with:

```
$ kubectl delete service cluster-example-rw
```

And get the list of services from Kubernetes:

```
$ kubectl get services
```

Why cluster-example-rw is still there? Is that the same service?

Try deleting cluster-example-2. What happened?

Try deleting cluster-example-1. What happened? Who is the primary now?

Promote manually cluster-example-2 with:

```
kubectl cnpg promote cluster-example 2
```



# Exercise #7 Configuring PostgreSQL

\$ kubectl get cluster cluster-example -o yaml |grep -A25 parameters:

parameters:

archive\_mode: "on"

archive\_timeout: 5min

dynamic\_shared\_memory\_type: posix

wal\_keep\_size: 512MB

wal\_receiver\_timeout: 5s

wal\_sender\_timeout: 5s

syncReplicaElectionConstraint:

enabled: false

primaryUpdateMethod: restart

primaryUpdateStrategy: unsupervised

1. Try setting ***log\_directory*** with kubectl edit cluster cluster-example. Is that possible? Why?
2. What is the current value of ***shared\_buffers***?
3. In a separate tab run kubectl get pod -w. Increase ***shared\_buffers***. What happened?
4. Decrease ***shared\_buffers***. What happened?



# Exercise #8 Scaling

Cluster-example has three instances (one primary and two replicas). Scale it to 4 with:

```
$ kubectl scale cluster cluster-example --replicas=4
```

Scale it back to 3

## Links

- [https://cloudnative-pg.io/documentation/current/operator\\_capability\\_levels/#scale-up-and-down-of-replicas](https://cloudnative-pg.io/documentation/current/operator_capability_levels/#scale-up-and-down-of-replicas)



# Exercise #9 Wreck it all

Remove the directory of the `postgres` database in the current primary:

```
kubectl exec cluster-example-<X> -- rm -rf /var/lib/postgresql/data/pgdata/base/5
```

What happened? Which Pod is the new primary?

```
$ kubectl exec -ti cluster-example-<Y> -- psql
```

Defaulted container "postgres" out of: postgres, bootstrap-controller (init)

psql (16.4 (Debian 16.4-1.pgdg110+2))

Type "help" for help.

```
postgres=# create table testtest (i integer);
```

```
CREATE TABLE
```

```
postgres=# insert into testtest (select generate_series(1,1000000));
```

```
INSERT 0 1000000
```

```
postgres=#
```

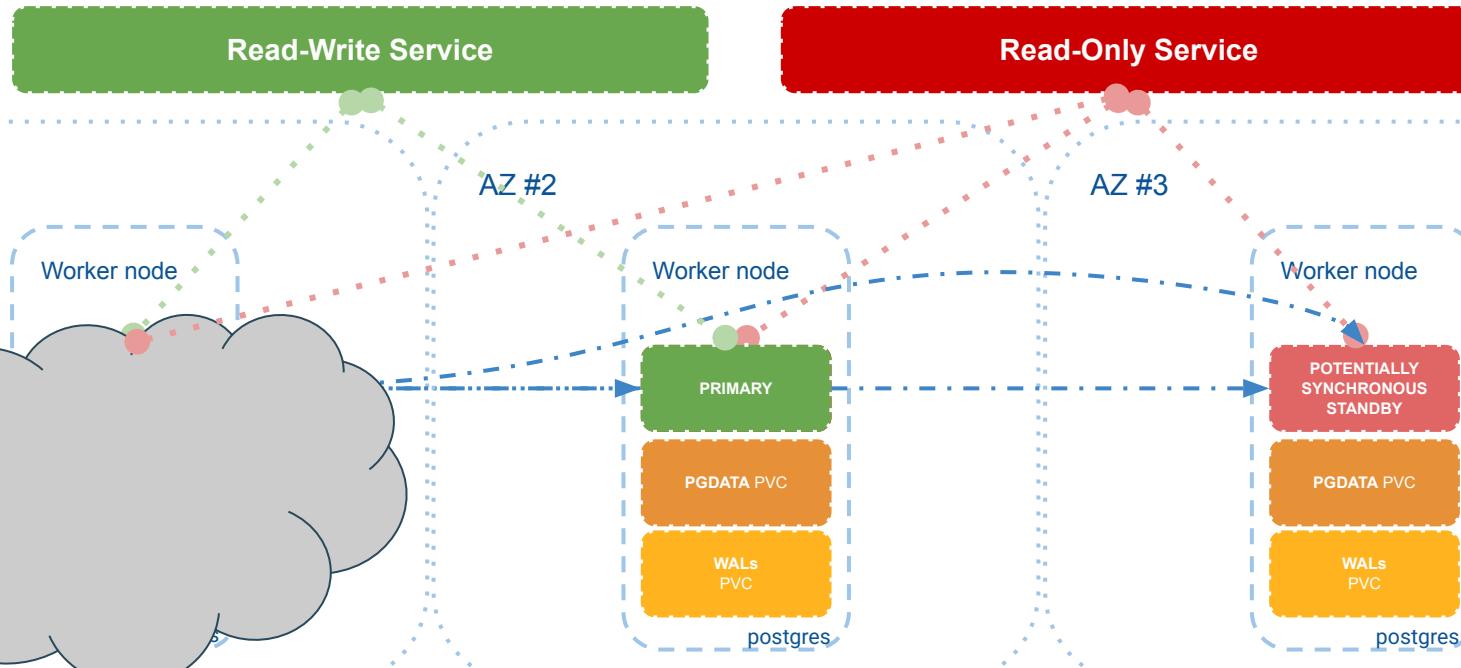
Destroy the Pod and its data with:

```
kubectl cnpq destroy cluster-example <X>
```



# Automated failover (HA with very low RTO)

K8s cluster



# Exercise #10 WAL Archiving

While connected to the EU region, apply the example in examples/eu/cluster-dc-eu.yaml:

```
$ kubectl apply -f examples/eu/cluster-dc-eu.yaml  
$ k get cluster cluster-dc-eu -o yaml|grep -A20 backup:
```

```
  backup:  
    barmanObjectStore:  
      destinationPath: s3://backups/  
      endpointURL: http://minio-eu:9000  
      s3Credentials:  
        accessKeyId:  
          key: ACCESS_KEY_ID  
          name: minio-eu  
        secretAccessKey:  
          key: ACCESS_SECRET_KEY  
          name: minio-eu
```

```
    wal:  
      compression: gzip  
      target: prefer-standby
```

```
  bootstrap:  
    initdb:  
      database: app  
      owner: app
```

Is it archiving WALs?

```
$ docker exec minio-eu ls /data/backups/cluster-dc-eu/wals/0000000100000000  
000000010000000000000001.gz
```



# Exercise #11 On Demand backups/PITR

```
$ kubectl cnpg backup cluster-dc-eu  
backup/cluster-dc-eu-20241007164438 created
```

```
$ kubectl get backup cluster-dc-eu-20241007164438  
NAME           AGE   CLUSTER      METHOD      PHASE    ERROR  
cluster-dc-eu-20241007164438  15s  cluster-dc-eu  barmanObjectStore  completed
```

```
$ kubectl get backup cluster-dc-eu-20241007164438 -o yaml
```

```
apiVersion: postgresql.cnpg.io/v1  
kind: Backup  
metadata:  
  creationTimestamp: "2024-10-07T14:44:38Z"  
  generation: 1  
  labels:  
    cnpg.io/cluster: cluster-dc-eu  
  name: cluster-dc-eu-20241007164438  
  namespace: default  
  resourceVersion: "4487"  
  uid: 61eb23dd-737f-4e19-879f-15f0487bebc1  
spec:  
  cluster:  
    name: cluster-dc-eu  
  method: barmanObjectStore
```



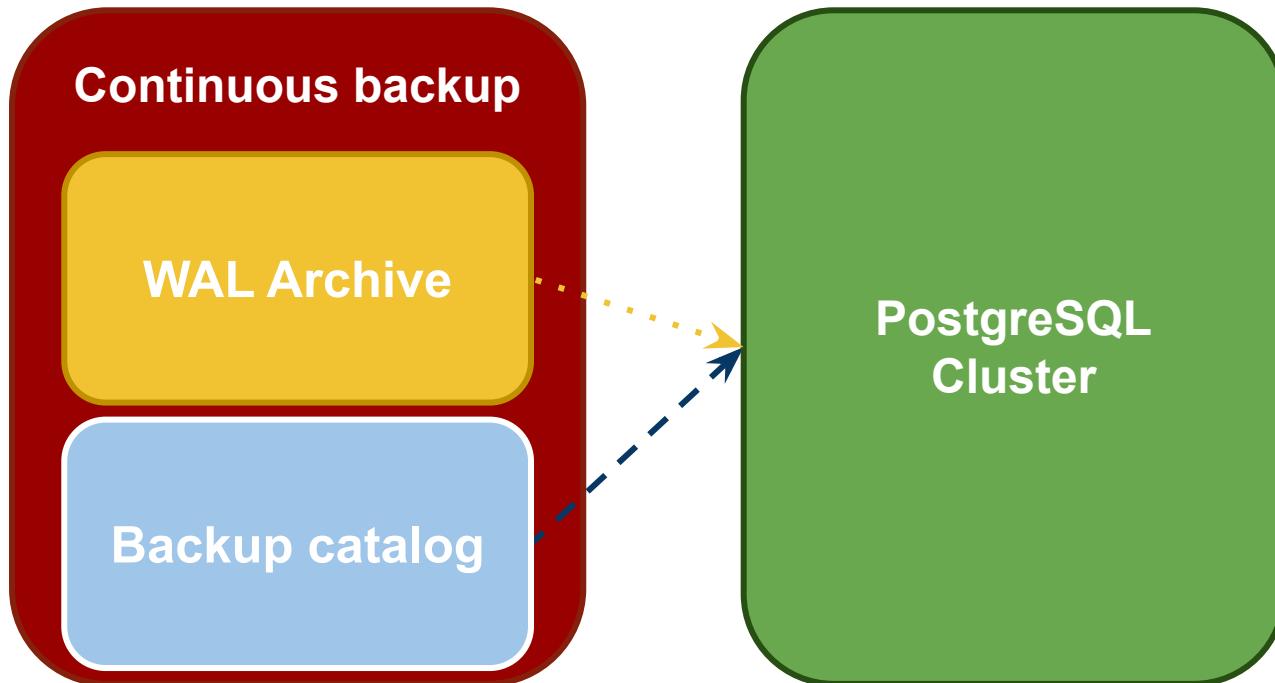
# Exercise #12 Creating new cluster from backup

```
$ kubectl apply -f - <<EOF
apiVersion: postgresql.cnpq.io/v1
kind: Cluster
metadata:
  name: restore
spec:
  instances: 3
  storage:
    size: 1Gi
  bootstrap:
    recovery:
      source: cluster-dc-eu
  externalClusters:
  - name: cluster-dc-eu
    barmanObjectStore:
      serverName: cluster-dc-eu
      destinationPath: s3://backups/
      endpointURL: http://minio-eu:9000
    s3Credentials:
      accessKeyId:
        name: minio-eu
        key: ACCESS_KEY_ID
      secretAccessKey:
        name: minio-eu
        key: ACCESS_SECRET_KEY
    wal:
      compression: gzip
EOF
```



# Full recovery or PITR

PostgreSQL is promoted when  
the recovery target is reached



# Exercise #13 Metrics

```
helm repo add prometheus-community \
https://prometheus-community.github.io/helm-charts
```

```
helm upgrade --install \
-f https://raw.githubusercontent.com/cloudnative-pg/cloudnative-pg/main/docs/src/samples/monitoring/kube-stack-config.yaml \
prometheus-community \
prometheus-community/kube-prometheus-stack
```

```
kubectl apply -f - <<EOF
---
apiVersion: postgresql.cnpq.io/v1
kind: Cluster
m
etadata:
  name: cluster-with-metrics
spec:
  instances: 3

  storage:
    size: 1Gi

  monitoring:
    enablePodMonitor: true
EOF
```

```
kubectl port-forward svc/prometheus-community-kube-prometheus 9090
kubectl port-forward svc/prometheus-community-grafana 3000:80
```



# Exercise #14 Hibernation and rehydration

## Create a basic cluster

```
apiVersion: postgresql.cnpq.io/v1
kind: Cluster
metadata:
  name: cluster-example
spec:
  imagePullPolicy: Always
  instances: 3
  storage:
    storageClass: standard
    size: 1Gi
```

```
kubectl annotate cluster cluster-example --overwrite cnpq.io/hibernation=on
```

```
kubectl get cluster cluster-example -o "jsonpath={.status.conditions[?(.type==\"cnpq.io/hibernation\")]}"
```

```
kubectl annotate cluster cluster-example --overwrite cnpq.io/hibernation=off
```



# Thank you

